

Unsupervised learning finds patterns in data, like clusters of customers by their purchases, or compressing the data using purchasing patterns (dimension reduction)

Unsupervised learning learns without labels, nor a specific prediction task in mind

Data will be written in 2D NumPy arrays. Columns will be features, and rows will be samples.

The samples of this dataset is in 4 dimensions.

Dimension = number of features

We can reduce dimensionality with k-means clustering from sklearn

```
from sklearn.cluster import KMeans
model = KMeans(n_clusters=3)
model.fit(samples)
labels = model.predict(samples)

# new data can be clustered without starting over, they remember the centroids
new_labels = model.predict(new_samples)

# we can use scatter plots with pyplot from matplotlib to visualize

import matplotlib.pyplot as plt
xs = samples[:,0]
ys = samples[:,2]
plt.scatter(xs, ys, c=labels) # color by cluster label
plt.show()
```

Also, we can use a different function to combine fit and predict

```
model.fit_predict()
```

We can put points and corresponding centroids on the same graph

```
# Import pyplot
import matplotlib.pyplot as plt

# Assign the columns of new_points: xs and ys
xs = new_points[:,0]
ys = new_points[:,1]

# Make a scatter plot of xs and ys, using labels to define the colors
plt.scatter(xs, ys, c=labels, alpha=0.5)
```

```
# Assign the cluster centers: centroids
centroids = model.cluster_centers_

# Assign the columns of centroids: centroids_x, centroids_y
centroids_x = centroids[:,0]
centroids_y = centroids[:,1]

# Make a scatter plot of centroids_x and centroids_y
plt.scatter(centroids_x, centroids_y, marker='D', s=50)
plt.show()
```

We can evaluate a customer by checking correspondence with the species

We can **cross tabulate** with pandas

```
import pandas as pd
df = pd.DataFrame({'labels': labels, 'species': species})

ct = pd.crosstab(df['labels'], df['species'])
```

We can measure clustering quality using only samples and their cluster labels

A good clustering has tight clusters, measured by the inertia, the distance from each sample to centroid of its cluster

```
# is it automatically measured
print(model.inertia_)
```

Sometimes, as number of clusters increases, inertia decreases

A good clustering has tight clusters (so low inertia), we need to choose an "elbow" point on the inertia plot for the right number of clusters

## Transforming features for better clusterings

Dataset: 178 samples, from 3 distinct varieties

Feature variance requires feature transformation, KMeans won't be effective

We can use StandardScaler to standardize features

```
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
scaler.fit(samples)
StandardScaler(copy=True, with_mean=True, with_std=True)
samples_scaled = scaler.transform(samples)
```

While KMeans fits and predicts, StandardScaler fits and scales

We can use a pipeline to combine these

```
from sklearn.pipeline import make_pipeline
pipeline = make_pipeline(scaler, kmeans)
pipeline.fit(samples)
labels = pipeline.predict(samples)
```